



trained to match the model's final output: $D_{\text{KL}}(P(\cdot | x) \parallel Q_{\ell, u})$